

MCP Security Checklist

A practical, community-maintained security baseline for teams building and deploying Model Context Protocol servers and AI agent infrastructure. Vendor-neutral. Free to use. Open to contributions.

6 SECURITY DOMAINS	36 TOTAL CONTROLS	2 CRITICAL CATEGORIES	3 AUDIENCE TIERS
---------------------------------	--------------------------------	------------------------------------	-------------------------------

QUICK START

Top 10 Controls

01 Never expose MCP over the public internet without mTLS or equivalent.	02 Scope every tool to the minimum necessary permissions.
03 Validate and sanitize all inputs before they reach tool execution.	04 Log every tool invocation with the originating session context.
05 Set rate limits on both the MCP server and any downstream APIs it calls.	06 Treat agent sessions as untrusted by default — validate intent, not just tokens.
07 Separate read and write tools; require explicit approval for write ops in sensitive contexts.	08 Rotate credentials used by MCP servers on a defined schedule.
09 Monitor for behavioral anomalies: unusual tool chains, high-frequency calls, off-hours access.	10 Conduct a tool inventory review before every production deployment.

Security Controls

Authentication & Authorization

CRITICAL

10 controls ▾

- Require authentication on all MCP endpoints**
Unauthenticated tool invocation allows any caller to execute tools.
auth-001
- Use short-lived tokens for agent sessions**
Long-lived static keys increase exposure window on compromise.
auth-002
- Implement mTLS for server-to-server MCP communication**
Mutual TLS ensures both parties are authenticated in transit.
auth-003
- Do not embed credentials in agent prompts or tool descriptions**
Agents can leak credentials in responses or logs.
auth-004
- Use a secrets manager for MCP server credentials**
Environment variables are frequently exposed in logs or config dumps.
auth-005
- Apply least-privilege scoping to every tool**
Overprivileged tools expand blast radius on compromise.
auth-006
- Separate read-only and write/execute tools**
Require explicit elevated authorization for destructive operations.
auth-007
- Enforce per-agent identity on tool invocations**
Shared credentials prevent individual agent traceability.
auth-008
- Validate tool-level authorization, not just server-level auth**
Gateway-only checks can be bypassed via direct MCP calls.
auth-009
- Restrict which tools are visible to which agents**
Tool discovery itself is an attack surface.
auth-010

Input Validation & Prompt Injection

CRITICAL

6 controls ▾

- Validate structure and type of all tool inputs before execution

Unvalidated inputs can cause unexpected behavior or injection.

input-001

- Define and enforce input schemas for every tool

Schema enforcement rejects malformed and adversarial inputs early.

input-002

- Set maximum length limits on all string inputs

Unbounded inputs can cause resource exhaustion.

input-003

- Separate data from instructions at the tool layer

Mixing data and instructions enables prompt injection via retrieved content.

input-004

- Validate output of tool results before passing back to model

Tool responses can contain adversarial instructions targeting further agent calls.

input-005

- Require explicit confirmation for irreversible actions

Reduces impact of successful prompt injection or misconfiguration.

input-006

Tool & Resource Exposure

HIGH

5 controls ▾

- Maintain a formal inventory of all exposed tools**
You cannot secure what you haven't catalogued.
tool-001
- Review tool inventory before every deployment**
Development tools frequently carry over-permissive access into production.
tool-002
- Expose only tools an agent needs for its current task**
Minimal tool surface limits blast radius of compromised sessions.
tool-003
- Use separate MCP instances for different agent roles**
Mixing admin and read-only tools enables privilege escalation.
tool-004
- Apply egress filtering on all external API calls from MCP tools**
Unfiltered egress enables data exfiltration via tool-initiated calls.
tool-005

API Session Security

HIGH

5 controls ▾

Assign a unique session ID to every agent interaction

Session IDs are required to trace multi-step tool chains.

session-001

Enforce session timeouts

Indefinite sessions increase the window for token abuse.

session-002

Apply rate limits per session and per tool

Server-level rate limiting alone is insufficient for agentic workloads.

session-003

Alert on sustained high-frequency tool calls from a single session

Indicates runaway loops, abuse, or compromised agents.

session-004

Log full request context for every tool call

Incomplete logs prevent incident reconstruction.

session-005

Monitoring & Observability

HIGH

5 controls ▾

- Log every tool invocation with session ID, agent identity, tool name, and timestamp
Minimum required data for incident response.

mon-001

- Log tool chain sequences within sessions
Unusual chains are a primary detection signal for agentic threats.

mon-002

- Centralize MCP logs into SIEM or log aggregation platform
Siloed logs are inaccessible during a server compromise.

mon-003

- Retain logs for a minimum of 90 days
Agentic attack chains can take time to fully manifest.

mon-004

- Build a behavioral baseline before relying on anomaly alerts
Without a baseline, anomaly detection generates excessive false positives.

mon-005

Network & Infrastructure

HIGH

5 controls ▾

Do not expose MCP servers directly to the public internet
Public exposure dramatically increases attack surface.
net-001

Place MCP servers behind a gateway or reverse proxy
Centralizes TLS termination, auth enforcement, and rate limiting.
net-002

Apply egress filtering to MCP servers
Prevents arbitrary outbound connections and exfiltration.
net-003

Run MCP server processes as a non-root, dedicated service account
Limits blast radius of OS-level compromise.
net-004

Keep MCP server dependencies up to date
Supply chain vulnerabilities in libraries are a real and exploited risk.
net-005

Maintained by [Helixar](https://helixar.ai) (https://helixar.ai)
security research · [GitHub](https://github.com/helixar-ai/mcp-security-checklist)
MIT License

[Contribute](https://github.com/helixar-ai/mcp-security-checklist/blob/main/CONTRIBUTING.md) (https://github.c
om/helixar-ai/mcp-security-
checklist/blob/main/CONTRIB
UTING.md)

helixar.ai
(https://
helixar.a
i)